

# PDANA8411 POE PART 3

Brice Derek Agnew (ST10072411)  
IIE VARSITY COLLEGE Cape Town

## Contents

Introduction:.....	2
The types of models used: .....	3
Plan: .....	4
Report:.....	5
References .....	16

## Introduction:

According to Marius Mogyorosim (Mogyorosi, 2025), Sentiment analysis in data analytics can be defined as the process of using Natural Language Processing and other methodologies to automatically determine the tone and emotion expressed in digital text(Mogyorosi, 2025).

In this project I will be making use of the NLTK library to perform various Sentiment analysis tasks in order to identify the prevalent sentiment in the dataset and identify any frequently occurring concerns in the reviews.

The dataset used in this project is the publicly available “**Business Reviews from different Industries**” dataset by user **Ashlin Darius Govindasamy** on Kaggle.com(Govindasamy, 2023). I chose it because it was a relatively large dataset with around 40 000 entries, consisting of ‘Hello Peter’ reviews, which is the same platform and thus format that the company would use.

## The types of models used:

This project made use of the **Sentiment Intensity Analyzer methodology**, an easily applicable sentiment analysis method from the 'nltk' library that doesn't require training and is easy to implement([www.nltk.com](http://www.nltk.com), 2025). The 'SIA' method is also very good at extracting sentiment based on punctations and emoticons, meaning that it could potentially be more accurate as it draws on more([www.nltk.com](http://www.nltk.com), 2025).

Logistic Regression, a classification algorithm that is well suited to the problem of sentiment analysis, it is however, better suited for labelled data and is not entirely required for the purpose of this project, I did however train one that can be used for 'Future sentiment prediction'. This additional model would prove to predict sentiment at an 86% accuracy score according to the analysis.

## Plan:

This project will follow the steps of the Data Analytics process, making amendments where necessary to suit the needs of the client as described in the question paper.

The first step in this process is to define the problem. Our problem for this project is 'How can we develop a sentiment analysis program, based on 'Hello Peter' reviews?'

The second step is data collection. I sourced my data as a csv file from a publicly available Kaggle dataset of hello peter reviews(Govindasamy, 2023), which aligns with our problem.

The third step is Exploratory Data Analysis, a means of visually representing and defining the types of columns in the dataset, number of missing values, number of rows, number of duplicate rows, number of missing values per column, and the datatypes of each column in the dataset.

The fourth step was Data Cleaning and Preprocessing, in which I made use of a pipeline to identify and remove any missing or duplicate values(pandas.pydata.org, 2025b), and remove any problematic columns(pandas.pydata.org, 2025a).

After data cleaning, Sentiment Analysis is undertaken using the Sentiment Intensity Analyzer package from the nltk library(www.nltk.com, 2025). This then assesses each review content to determine if the sentiment is positive, negative, or neutral. This in turn creates a label for each data entry.

Next, a Logistic regression model is trained to predict future sentiments, if necessary, for future implementation.

After the model is developed The final stage of the project is 'Concern Identification', in which I used the nltk 'tokenize'(www.nltk.org, n.d.) and 'corpus'(www.nltk.org, n.d.) packages to tokenize all the review data in the dataset, sort through them to determine the total frequency of each prominent word used in the reviews, assign each of the top 50 most frequently occurring words into 7 key categories, and determining and displaying the total frequency of each 'concern category' to determine which primary concerns were discussed in the reviews. According to the data analysis I undertook, the most frequently discussed concern was "Customer Service" with 127735 observed mentions in the dataset.

## Report:

The preliminary step of this project is to ensure that the nltk package is installed on the local machine:

```
[1]: !pip install nltk

Requirement already satisfied: nltk in c:\users\brice\anaconda3\lib\site-packages (3.9.1)
Requirement already satisfied: click in c:\users\brice\anaconda3\lib\site-packages (from nltk) (8.1.7)
Requirement already satisfied: joblib in c:\users\brice\anaconda3\lib\site-packages (from nltk) (1.4.2)
Requirement already satisfied: regex>=2021.8.3 in c:\users\brice\anaconda3\lib\site-packages (from nltk) (2024.9.11)
Requirement already satisfied: tqdm in c:\users\brice\anaconda3\lib\site-packages (from nltk) (4.66.5)
Requirement already satisfied: colorama in c:\users\brice\anaconda3\lib\site-packages (from click->nltk) (0.4.6)
```

The relevant imports are added to the program to ensure access to relevant methods, functions, and models throughout the project:

```
[2]: #ST10072411 Brice Agnew
#24/06/2025
#Prog8411 POE part 3
#Sentiment Analysis

#Imports
import nltk

try:
    nltk.data.find('sentiment/vader_lexicon')
except LookupError:
    nltk.download('vader_lexicon')
# tries to find the vader_lexicon, and if not, downloads to the 'C:\Users\<user>\AppData\Roaming\nltk_data' directory

from nltk.sentiment.vader import SentimentIntensityAnalyzer as SIA
import pandas as pd
import re

from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import SelectKBest, chi2
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from sklearn.preprocessing import LabelEncoder

from sklearn.base import BaseEstimator, TransformerMixin
from wordcloud import WordCloud
import matplotlib.pyplot as plt

from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
from collections import Counter

# Download only once, ensuring that required NLTK resources are available
nltk.download('punkt')
nltk.download('stopwords')
```

After this the raw data file(Govindasamy, 2023) is imported and sved as a dataframe for use in the project:

```
•[3]: #Data Collection

#open chosen dataset csv file
#Please be sure to keep the attached csv files in the same root repository before running!
Dataset = pd.read_csv("PDANA_POE_rawdata.csv")
```

Next we can begin the Exploratory Data Analysis stage, in order to identify key aspects of the dataset:

[4]:

#Exploratory Data Analysis: Part 1

Dataset.head(7)

[4]:		id	user_id	created_at	authorDisplayName	author	authorAvatar	author_id	review_title	review_rating	review_content	...	industry_name
	0	4228893	0cf37140-f94c-11ec-9b49-19cd672a8c8b	2022-12-16 20:51:27	Werner S	Werner S	NaN	0cf37140-f94c-11ec-9b49-19cd672a8c8b	Not telkpm	1	I pay R519 per month and can't signon the ...	...	Telecommunication:
	1	4228846	364faaf0-5104-11ec-8045-7d7a9af8c558	2022-12-16 18:54:28	Nokuthula K	Nokuthula K	NaN	364faaf0-5104-11ec-8045-7d7a9af8c558	Please help me clear my name	1	I can't even buy a house because telkom failed...	...	Telecommunication:
	2	4228766	0c8502a0-5ee0-11ea-9899-6f17cff30ea1	2022-12-16 16:28:46	Services A	Services A	NaN	0c8502a0-5ee0-11ea-9899-6f17cff30ea1	Untrustworthy business	1	I was contacted if I wanted to upgrade a data ...	...	Telecommunication:
	3	4228626	35bf64d7-31fa-11e8-83f4-f23c91bb6188	2022-12-16 12:58:01	Nina H	Nina H	NaN	35bf64d7-31fa-11e8-83f4-f23c91bb6188	telkom	1	Have been in a fight with TELKOM since August ...	...	Telecommunication:
	4	4228610	2b490ce9-31fa-11e8-83f4-f23c91bb6188	2022-12-16 12:42:24	Arshad P	Arshad P	NaN	2b490ce9-31fa-11e8-83f4-f23c91bb6188	TELKOM Killarney Looting on Upgrade	1	-10, \nSo I have Telkom contract with Telkom f...	...	Telecommunication:
	5	4228533	017368a3-31fa-11e8-83f4-f23c91bb6188	2022-12-16 11:27:08	Elaine	Elaine	NaN	017368a3-31fa-11e8-83f4-f23c91bb6188	TELKOM AND NUDEBT ARE ROGUES AND FRAUDSTERS!!!	1	TELKOM ARE ROGUES, \nTHEY STEAL OUR MONEY EVER...	...	Telecommunication:
	6	4228398	3c07bdd1-31fa-11e8-83f4-f23c91bb6188	2022-12-16 08:48:06	Robyn S	Robyn S	NaN	3c07bdd1-31fa-11e8-83f4-f23c91bb6188	Despicable Service	1	If I could give this company a 1 I would. The ...	...	Telecommunication:

7 rows × 27 columns

In this stage of EDA, the goal is to visually represent the properties of the dataset through visual or tabular formats to convey information. The following visual is used to proportionately represent the distribution of most frequently occurring words in the dataset:

[illegible]

The following code identifies key properties of the dataset before cleaning, such as the number of rows, the number of duplicate values, the number of missing values in the set, and the datatypes used for each column in the raw dataset:

```
[6]: #Exploratory Data Analysis: Part 3
print("Dataset Empty value check: ")
print(DataSet.isna())
print("")

print("Number of rows:\n"
      , len(DataSet), "\n")
print("Number of Empty values:\n"
      , DataSet.isna().any(axis=1).sum(), "\n")

#Identifying number of duplicate rows
print("Number of Duplicate rows:\n"
      , DataSet.duplicated().sum(), "\n")

#Counting which columns have the most missing values
print("Number of missing values per column: ")
print(DataSet.isna().sum())
print("")

#investigating datatypes and column names
print("Data types per column: ")
print(DataSet.dtypes)
```

```
Dataset Empty value check:
   id  user_id  created_at  authorDisplayName  author  authorAvatar \
0  False  False      False                False  False      True
1  False  False      False                False  False      True
2  False  False      False                False  False      True
3  False  False      False                False  False      True
4  False  False      False                False  False      True
...    ...    ...          ...              ...    ...      ...
56255  False  False      False                False  False      True
56256  False  False      False                False  False      True
56257  False  False      False                False  False      True
56258  False  False      False                False  False      True
56259  False  False      False                False  False      True

   author_id  review_title  review_rating  review_content  ... \
0  False      False      False      False      False  ...
1  False      False      False      False      False  ...
2  False      False      False      False      False  ...
3  False      False      False      False      False  ...
4  False      False      False      False      False  ...
...    ...    ...          ...              ...    ...  ...
56255  False      False      False      False      False  ...
56256  False      False      False      False      False  ...
56257  False      False      False      False      False  ...
56258  False      False      False      False      False  ...
56259  False      False      False      False      False  ...

   industry_name  industry_slug  status_id  nps_rating  source \
0  False      False      False      True      False
1  False      False      False      True      False
2  False      False      False      True      False
3  False      False      False      True      False
4  False      False      False      True      False
...    ...    ...          ...              ...    ...
56255  False      False      False      True      True
56256  False      False      False      True      True
56257  False      False      False      True      True
56258  False      False      False      True      True
56259  False      False      False      True      True

   is_reported  business_reporting  author_created_date \
0  False      True                False
1  False      True                False
2  False      True                False
3  False      True                False
4  False      True                False
...    ...    ...                ...
56255  False      True                False
56256  False      True                False
56257  False      True                False
56258  False      True                False
56259  False      True                False

   author_total_reviews_count  attachments
0  False                    False
1  False                    False
2  False                    False
3  False                    False
4  False                    False
...    ...                    ...
56255  False                    False
56256  False                    False
56257  False                    False
56258  False                    False
56259  False                    False
```

[56260 rows x 27 columns]



```

Number of rows:
56260

Number of Empty values:
56259

Number of Duplicate rows:
0

Number of missing values per column:
id                0
user_id           5
created_at        0
authorDisplayName 0
author            0
authorAvatar      55976
author_id         5
review_title      6
review_rating     0
review_content    14
business_name     0
business_slug     0
permalink         0
replied          0
messages          0
business_logo     11086
industry_logo     0
industry_name     0
industry_slug     0
status_id         0
nps_rating        40745
source            11319
is_reported       0
business_reporting 55700
author_created_date 0
author_total_reviews_count 0
attachments       0
dtype: int64

Data types per column:
id                int64
user_id           object
created_at        object
authorDisplayName object
author            object
authorAvatar      object
author_id         object
review_title      object
review_rating     int64
review_content    object
business_name     object
business_slug     object
permalink         object
replied          int64
messages          object
business_logo     object
industry_logo     object
industry_name     object
industry_slug     object
status_id         int64
nps_rating        float64
source            object
is_reported       bool
business_reporting object
author_created_date object
author_total_reviews_count int64
attachments       object
dtype: object

```

The following code defines the key functions that will be used in the pipeline that will be used later. These functions define how the data will be cleaned and the functions selected:

```
[*]: #Data Cleaning Pipeline:
class ReviewCleaner(BaseEstimator, TransformerMixin):
    def __init__(self):
        # Define all irrelevant columns here for both Feature selection and Data Cleaning Steps
        self.cols_to_drop = [
            # Previously cleaned ones
            'business_reporting', 'authorAvatar', 'nps_rating', 'source', 'business_logo',
            # Feature selection irrelevant ones
            'id', 'user_id', 'created_at', 'messages', 'authorDisplayName', 'author',
            'author_id', 'attachments', 'industry_logo', 'industry_name', 'industry_slug',
            'is_reported', 'status_id', 'replied', 'author_created_date', 'permalink',
            'business_slug', 'business_name'
        ]

    def clean_text(self, text):
        # Remove all characters except letters, numbers, spaces, and punctuation
        return re.sub(r'^a-zA-Z0-9\s!?.]', '', text)

    def fit(self, X, y=None):
        return self

    def transform(self, X, y=None):
        X = X.copy()

        # Drop irrelevant columns if they exist
        X.drop(columns=[col for col in self.cols_to_drop if col in X.columns], inplace=True, errors='ignore')

        # Combine and clean review_title + review_content (For sentiment analysis and Theme identification)
        if {'review_title', 'review_content'}.issubset(X.columns):
            X['review'] = (X['review_title'].fillna('') + ' ' + X['review_content'].fillna(''))
            X['review'] = X['review'].apply(self.clean_text)
            X = X.drop(columns=['review_title', 'review_content'], errors='ignore')

            X = X.drop_duplicates()
        return X

class TextSelector(BaseEstimator, TransformerMixin):
    def __init__(self, key='review'):
        self.key = key

    def fit(self, X, y=None):
        return self

    def transform(self, X):
        return X[self.key] # return a Series, not DataFrame
```

This code shows how the pipeline is implemented:

```
[8]: #Data Cleaning: Part 2
ml_pipeline = Pipeline([
    ('cleaner', ReviewCleaner()), # expects a DataFrame
    ('select_text', TextSelector('review')), # reduce to just the 'review' column
    ('vectoriser', TfidfVectorizer(stop_words='english')),
    ('select', SelectKBest(chi2, k=1000)),
    ('model', LogisticRegression(max_iter=1000))
])
```

The following is the code and output of the usage of the pipeline, the dataset is cleaned and transformed into the 'cleaned\_data' dataframe, with a test output and expected output to confirm correct data transformation:

```
[9]: #Data Cleaning: Part 3
#Pipeline call
# Apply only the cleaning step
cleaned_data = ReviewCleaner().fit_transform(Dataset)

#Establishing the number of empty values remaining in the Dataset
print("Number of Empty values:\n",
      cleaned_data.isna().any(axis=1).sum(), "\n")

print(cleaned_data.isna().sum())
print("")

#Testing the new 'review' content
print("The review content: \n", cleaned_data['review'][0])
#Expected output: 'Not telkpm I pay R519 per month and can't signin on the modem. This happens when we have electricity which is about every 3 hours none after 21h00
#I would plead with potential customers not to consider telkom'
```

Number of Empty values:  
0

review_rating	0
author_total_reviews_count	0
review	0

dtype: int64

The review content:  
Not telkpm I pay R519 per month and cant signin on the modem. This happens when we have electricity which is about every 3 hours none after 21h00  
I would plead with potential customers not to consider telkom

The next section of code demonstrates how the data has been transformed by displaying the number of rows, empty values, duplicate values, and columns that are present in the cleaned dataset, in comparison to the raw dataset:

```
[10]: #Data Cleaning: Part 4
#visually representing changes
print("Number of rows:\n",
      len(cleaned_data), "\n")
print("Number of Empty values:\n",
      cleaned_data.isna().any(axis=1).sum(), "\n")

#Identifying number of duplicate rows
print("Number of Duplicate rows:\n",
      cleaned_data.duplicated().sum(), "\n")

#Counting which columns have the most missing values
print(cleaned_data.isna().sum())
print("")
```

Number of rows:  
55926

Number of Empty values:  
0

Number of Duplicate rows:  
0

review_rating	0
author_total_reviews_count	0
review	0

dtype: int64

This following section of code shows how the Sentiment Intensity Analyzer module is used in the program to assign levels of sentiment to each row, as well as the columns that are created in this process:

```
[11]: #Sentiment Analysis:
#using the SentimentIntensityAnalyzer package from the NLTK library, specifically the vader
#SIA doesnt require training, which does make the process simpler. Additionally it is reported to be good at interpreting punctuation in sentiment.

#calculate sentiment
analyzer = SIA()
cleaned_data['compound'] = [analyzer.polarity_scores(x)['compound'] for x in cleaned_data['review']]
cleaned_data['negative'] = [analyzer.polarity_scores(x)['neg'] for x in cleaned_data['review']]
cleaned_data['neutral'] = [analyzer.polarity_scores(x)['neu'] for x in cleaned_data['review']]
cleaned_data['positive'] = [analyzer.polarity_scores(x)['pos'] for x in cleaned_data['review']]

#create columns
cleaned_data['sentiment']='neutral' #Automatically setting 'overall' sentiment to being neutral before individually assigning positive or negative
cleaned_data.loc[cleaned_data.compound>0.05,'sentiment']='positive'
cleaned_data.loc[cleaned_data.compound<-0.05,'sentiment']='negative'
cleaned_data.head()

#on-screen summary
print("Neutral Total Sentiment Count: ",
      cleaned_data['sentiment'].value_counts()['neutral'])
print("Positive Total Sentiment Count: ",
      cleaned_data['sentiment'].value_counts()['positive'])
print("Negative Total Sentiment Count: ",
      cleaned_data['sentiment'].value_counts()['negative'])

cleaned_data.head(5)
```

Neutral Total Sentiment Count: 1884  
Positive Total Sentiment Count: 34192  
Negative Total Sentiment Count: 19850

```
[11]:
```

	review_rating	author_total_reviews_count	review	compound	negative	neutral	positive	sentiment
0	1	3	Not telkpm I pay R519 per month and cant signi...	0.0762	0.000	0.963	0.037	positive
1	1	3	Please help me clear my name I cant even buy a...	0.8608	0.055	0.768	0.177	positive
2	1	50	Untrustworthy business I was contacted if I wa...	-0.5994	0.079	0.890	0.031	negative
3	1	1	telkom Have been in a fight with TELKOM since ...	-0.9672	0.216	0.706	0.078	negative
4	1	5	TELKOM Killarney Looting on Upgrade 10 \nSo I ...	-0.9656	0.101	0.861	0.038	negative

The following code shows how a Logistic regression model is trained in order to predict the sentiment labels that were previously defined in the last sections, the model was then evaluated based on the metrics of Precision, Recall, f1-scores, and support.

```

•[12]: #Model training:
required_columns = {'review', 'sentiment'}

# Check if required columns exist
if not required_columns.issubset(cleaned_data.columns):
    missing = required_columns - set(cleaned_data.columns)
    raise ValueError(f"Columns missing from dataset: {missing}")

# Proceed only if columns exist
cleaned_data = cleaned_data.dropna(subset=['review', 'sentiment'])
#cleaned_data= cleaned_data.drop(columns=['review_title', 'review_content'], errors='ignore')

X = cleaned_data # needed by ReviewCleaner
y = cleaned_data['sentiment']

# Convert labels to numeric if needed
le = LabelEncoder()
y_encoded = le.fit_transform(y)

print("Length of X:", len(X))
print("Length of y:", len(y_encoded))

# Split data
#X_train, X_test, y_train, y_test = train_test_split(X, y_encoded, test_size=0.2, random_state=69)
X_train, X_test, y_train, y_test = train_test_split(X, y_encoded, test_size=0.2, random_state=69)

# Fit model
ml_pipeline.fit(X_train, y_train)

# Predict
preds = ml_pipeline.predict(X_test)

# Evaluate
print(classification_report(y_test, preds, target_names=le.classes_))

cleaned_data.head()

```

```

Length of X: 55926
Length of y: 55926

```

	precision	recall	f1-score	support
negative	0.80	0.90	0.85	3934
neutral	0.82	0.12	0.21	354
positive	0.91	0.89	0.90	6898
accuracy			0.87	11186
macro avg	0.85	0.63	0.65	11186
weighted avg	0.87	0.87	0.86	11186

```

[12]:

```

	review_rating	author_total_reviews_count	review	compound	negative	neutral	positive	sentiment
0	1	3	Not telkpm I pay R519 per month and cant signi...	0.0762	0.000	0.963	0.037	positive
1	1	3	Please help me clear my name I cant even buy a...	0.8608	0.055	0.768	0.177	positive
2	1	50	Untrustworthy business I was contacted if I wa...	-0.5994	0.079	0.890	0.031	negative
3	1	1	telkom Have been in a fight with TELKOM since ...	-0.9672	0.216	0.706	0.078	negative
4	1	5	TELKOM Killarney Looting on Upgrade 10 \nSo I ...	-0.9656	0.101	0.861	0.038	negative

The following code section is used to develop a visual representation of word distribution in the cleaned dataset:

```
[13]: #Visualization of word frequency in reviews
text = ''.join(cleaned_data['review'].dropna())
wordcloud = WordCloud(width=800, height=400, background_color='white', stopwords='english').generate(text)

plt.figure(figsize=(15, 7))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title("Most Common Words in Reviews")
plt.show()
```



The following code is used to tokenise every word used in the reviews in the dataset, and rank them based on most frequently used, ranking the top 50 most used words:

```

•[14]: #Concern Identification: Part 1
#Most frequently occurring words/concerns

# Prepare stopwords as a set for faster lookup
stop_words = set(stopwords.words('english'))

# Initialize counter
token_counts = Counter()

# Iterate over each review
for review in cleaned_data['review'].dropna().astype(str):
    tokens = word_tokenize(review.lower())
    tokens = [t for t in tokens if t.isalpha() and t not in stop_words]
    token_counts.update(tokens)

#the first version I tried with this involved compiling every review into a string to search through, but that had insane levels of overhead
#so this secondary version was developed to instead use a counter object and for loop to more efficiently search through the tokens found.

#Uses the token_count counter data to create a dataframe for display.
freq_df = pd.DataFrame(token_counts.items(), columns=['word', 'frequency'])
freq_df['frequency'] = freq_df['frequency'].astype(int)
freq_df = freq_df.sort_values(by='frequency', ascending=False)

print("Most frequently occurring words:")
freq_df.head(50)

```

Most frequently occurring words:

```

[14]:
      word  frequency
127  service    43802
153   call    16070
 70    get    15030
 12   would    12597
219   time    12015
 61   still    11223
499 received    11032
162   told    10778
1975  thank    10677
392   back    10490
 81  account    10152
420    one    10017
 17  telkom     9941
1180  great     9428
 22   even     9258
209  customer    9115
1693   car     8903
221  called     8822
384   good     8648
5345   br     8233
 1    pay     7933
195  money     7897
5747  excellent    7833
255   dont     7832

```

This final section of code takes the data insights extracted from the previous section and manually assigns each of the most frequently used top fifty words into different categories of concerns, with it then counted the occurrences of to determine the ranking of discussing themes in the reviews overall:

```
[15]: #Concern Identification: Part 2
#Manually seperating concerns into themes, and identifying the most commonly occuring themes

themes = {
    'Billing': ['pay', 'money', 'account', 'paid', 'bank', 'number', 'debit', 'bill'],
    'Customer Service': ['call', 'called', 'told', 'said', 'email', 'customer', 'help', 'please', 'service', 'sent', 'br'],
    'Product': ['telkom', 'company', 'network', 'phone', 'work', 'internet'],
    'Timeliness': ['time', 'days', 'day', 'month', 'still', 'back', 'never'],
    'Complaints': ['dont', 'need', 'want', 'im', 'never', 'get', 'like', 'made', 'one'],
    'Positive Sentiment': ['thank', 'great', 'good', 'excellent'],
    'Ambiguous': ['br', 'p']
}

theme_counts = {}

#iterating through the previous frequency distribution to assign each keyword to a theme
for theme, keywords in themes.items():
    total = sum(freq_df[freq_df['word'].isin(keywords)]['frequency'])
    theme_counts[theme] = total

theme_df = pd.DataFrame(list(theme_counts.items()), columns=['Concern', 'Mentions'])

#Presenting the most commonly occurring
print("[Most commonly occurring concerns:]")

theme_df.head(20).sort_values(by='Mentions', ascending=False)

[Most commonly occurring concerns:]
```

	Concern	Mentions
1	Customer Service	127479
4	Complaints	71285
3	Timeliness	60504
0	Billing	46656
5	Positive Sentiment	36586
2	Product	31441
6	Ambiguous	13674

The final section in the project contains references to that were used in informing the creation of this project:

```
[16]: #References:

#Govindasamy, A.D., 2023. Business Reviews from different Industries. [online] kaggle.com. Available at: <https://www.kaggle.com/datasets/ashlingovindas/
#Mogyorosi, M., 2025. Sentiment Analysis: First Steps With Python's NLTK Library. [online] realpython.com. Available at: <https://realpython.com/python-r
#pandas.pydata.org, 2025a. pandas.DataFrame.drop. [online] pandas.pydata.org. Available at: <https://pandas.pydata.org/pandas-docs/stable/reference/api/p
#pandas.pydata.org, 2025b. pandas.DataFrame.dropna. [online] pandas.pydata.org. Available at: <https://pandas.pydata.org/pandas-docs/stable/reference/api
#www.nltk.com, 2025. nltk.sentiment.SentimentIntensityAnalyzer. [online] www.nltk.com. Available at: <https://www.nltk.org/api/nltk.sentiment.SentimentIr
#www.nltk.org, n.d. nltk.corpus package. www.nltk.org. [online] Available at: <https://www.nltk.org/api/nltk.corpus.html> [Accessed 27 June 2025a].
#www.nltk.org, n.d. nltk.tokenize package. www.nltk.org. [online] Available at: <https://www.nltk.org/api/nltk.tokenize.html> [Accessed 27 June 2025b].
```



## References

- Govindasamy, A.D., 2023. *Business Reviews from different Industries*. [online] kaggle.com. Available at: <<https://www.kaggle.com/datasets/ashlingovindasamy/business-reviews-from-different-industries?select=trainingData.csv>> [Accessed 24 June 2025].
- Mogyorosi, M., 2025. *Sentiment Analysis: First Steps With Python's NLTK Library*. [online] realpython.com. Available at: <<https://realpython.com/python-nltk-sentiment-analysis/>> [Accessed 24 June 2025].
- pandas.pydata.org, 2025a. *pandas.DataFrame.drop*. [online] pandas.pydata.org. Available at: <<https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.drop.html>> [Accessed 24 June 2025].
- pandas.pydata.org, 2025b. *pandas.DataFrame.dropna*. [online] pandas.pydata.org. Available at: <<https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.dropna.html>> [Accessed 24 June 2025].
- www.nltk.com, 2025. *nltk.sentiment.SentimentIntensityAnalyzer*. [online] www.nltk.com. Available at: <<https://www.nltk.org/api/nltk.sentiment.SentimentIntensityAnalyzer.html?highlight=sentimentintensity>> [Accessed 27 June 2025].
- www.nltk.org, n.d. nltk.corpus package. *www.nltk.org*. [online] Available at: <<https://www.nltk.org/api/nltk.corpus.html>> [Accessed 27 June 2025a].
- www.nltk.org, n.d. nltk.tokenize package. *www.nltk.org*. [online] Available at: <<https://www.nltk.org/api/nltk.tokenize.html>> [Accessed 27 June 2025b].